

# Experiment Report on SIFT Algorithm Implementation

Wu Han  
Class: SEIEE 2301

December 8, 2024

## 1 Experiment Overview

This experiment provides an introduction to the SIFT (Scale-Invariant Feature Transform) algorithm, which was proposed by David Lowe in 1999 and further refined in 2004. The SIFT algorithm is used to detect and describe local features in images, making it robust against transformations such as translation, rotation, and affine changes. The SIFT algorithm has many key properties, like invariance, distinctiveness, multiplicity and scalability.

In OpenCV, there are pre-built functions for feature extraction and matching using SIFT. However, in this experiment, we attempt to implement our own version of a SIFT class to understand the underlying process.

The SIFT function implementation in this experiment follows these main steps:

1. Keypoint detection and localization.
2. Image preprocessing with Gaussian blur, followed by gradient magnitude and orientation computation.
3. Computation of the dominant orientation for each keypoint.
4. SIFT descriptor calculation for each keypoint.
5. Feature matching and visualization of results.

## 2 Keypoint Detection and Localization

For keypoint detection, we use the Harris corner detection method, which is implemented as follows:

```
1 def build_image_pyramid(image, levels=3):  
2     """  
3     Construct an image pyramid by resizing the image to multiple  
4     scales.  
5     """  
6     pyramid = [image]  
7     for _ in range(1, levels):
```

```

7     image = cv2.pyrDown(image)  # Downscale the image
8     pyramid.append(image)
9     return pyramid
10
11 def compute_harris_corners(image, block_size=2, ksize=3, k=0.04):
12     """
13     Use Harris corner detection to find keypoints with an
14     adaptive threshold.
15     """
16     corners = cv2.cornerHarris(image, blockSize=block_size, ksize
17                               =ksize, k=k)
18     corners = cv2.dilate(corners, None)  # Enhance corner points
19     adaptive_thresh = 0.25 * corners.max()  # Compute adaptive
20     threshold based on maximum response
21     keypoints = np.argwhere(corners > adaptive_thresh)
22     keypoints = [cv2.KeyPoint(float(p[1]), float(p[0]), 1) for p
23                 in keypoints]
24     return keypoints

```

The first function generates an image pyramid, which includes the original image and scaled-down versions. The pyramid helps detect features at multiple scales. **Key Points:**

- `cv2.pyrDown(image)` reduces the resolution by half in each step.
- The function returns a list of images, including the original and scaled-down versions.

The second function uses the Harris corner detection method to identify keypoints. An adaptive threshold selects strong corners.

### 3 Gaussian Blur and Gradient Calculation

The image is processed with Gaussian blur to smooth it. Then, the gradient magnitude and orientation are computed using the following formulas:

$$m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y-1) - L(x, y+1))^2}$$

$$\theta(x, y) = \arctan\left(\frac{L(x+1, y) - L(x-1, y)}{L(x, y-1) - L(x, y+1)}\right)$$

The gradients are computed in both directions, and any negative angles are adjusted to the range of  $[0, 2\pi]$ .

```

1 descriptors = []
2 image = cv2.GaussianBlur(image, (5, 5), 1.6)  # Gaussian
3         smoothing to reduce noise.
4
5 for kp in keypoints:
6     x, y = int(kp.pt[0]), int(kp.pt[1])  # Keypoint
7         coordinates.
8
9     # Extract the 16x16 neighborhood around the keypoint.

```

```

8     patch = image[max(0, y - patch_size // 2):y + patch_size
9                  // 2,
10                  max(0, x - patch_size // 2):x + patch_size
11                  // 2]
12
13     # Skip incomplete patches near image boundaries.
14     if patch.shape[0] != patch_size or patch.shape[1] !=
15        patch_size:
16         continue
17
18     # Compute gradients using Sobel operators.
19     gx = cv2.Sobel(patch, cv2.CV_32F, 1, 0, ksize=3)
20     gy = cv2.Sobel(patch, cv2.CV_32F, 0, 1, ksize=3)
21     magnitude, angle = cv2.cartToPolar(gx, gy, angleInDegrees
22                                       =True)

```

- The image is smoothed using a Gaussian filter with `cv2.GaussianBlur` to reduce noise.
- Gradients are calculated for the 16x16 neighborhood around each keypoint using Sobel operators:
  - `gx` represents the horizontal gradient.
  - `gy` represents the vertical gradient.
- Gradient magnitudes and orientations are computed using `cv2.cartToPolar`.

## 4 Dominant Orientation Calculation

To calculate the dominant orientation of each keypoint, a region around each keypoint is sampled. The region is weighted based on the distance to the center, and a histogram of gradients is computed. The dominant orientation is the one with the highest value in the histogram.

Each keypoint's descriptor is computed by creating a 16x16 block around the keypoint. This block is divided into 16 smaller 4x4 regions. For each region, bilinear interpolation is used to distribute gradient magnitudes across bins and cells. The descriptor is a 128-dimensional vector formed by the weighted sum of gradient magnitudes in each direction.

```

1 # Adjust angles relative to the dominant keypoint orientation.
2     main_orientation = kp.angle if kp.angle else 0
3     adjusted_angle = (angle - main_orientation) % 360

```

- The dominant orientation of the keypoint is calculated from the gradient data.
- Angles are adjusted relative to this orientation to ensure rotational invariance.

```

1 # Initialize a 4x4 grid with 8 bins per cell (128-dimensional
2 descriptor).
3     cell_size = patch_size // 4 # Each cell is 4x4 pixels.

```

```

3      descriptor = np.zeros((4, 4, 8), dtype=np.float32)
4
5      # Iterate over all pixels in the 16x16 patch.
6      for i in range(patch_size):
7          for j in range(patch_size):
8              # Compute relative coordinates within the grid.
9              patch_x = j + 0.5 # Pixel center adjustment
10             patch_y = i + 0.5
11
12             cell_x = patch_x / cell_size # Grid row index
13             cell_y = patch_y / cell_size # Grid column index
14
15             # Identify the 4 nearest grid cells.
16             x0, y0 = int(cell_x), int(cell_y)
17             x1, y1 = min(x0 + 1, 3), min(y0 + 1, 3) # Ensure
18                 within 4x4 grid boundaries.
19
20             # Compute interpolation weights.
21             dx1, dy1 = cell_x - x0, cell_y - y0
22             dx2, dy2 = 1 - dx1, 1 - dy1
23
24             # Quantize angle into 8 bins (each 45 degrees).
25             bin_idx = int(adjusted_angle[i, j] // 45) % 8
26             magnitude_value = magnitude[i, j]
27
28             # Bilinearly interpolate the gradient magnitude
29                 into the 4 nearest cells.
30             descriptor[y0, x0, bin_idx] += magnitude_value *
31                 dx2 * dy2
32             descriptor[y0, x1, bin_idx] += magnitude_value *
33                 dx1 * dy2
34             descriptor[y1, x0, bin_idx] += magnitude_value *
35                 dx2 * dy1
36             descriptor[y1, x1, bin_idx] += magnitude_value *
37                 dx1 * dy1

```

- The 16x16 patch is divided into a 4x4 grid, with each cell being 4x4 pixels.
- For each cell, an 8-bin histogram is generated to represent gradient orientations.
- Bilinear interpolation is used to distribute gradient magnitudes across bins and cells.
- The resulting histograms are concatenated to form a 128-dimensional descriptor.
- The descriptor is normalized to unit length, clipped to 0.2 to suppress extreme values, and re-normalized.

## 5 Feature Matching

To match descriptors between two images, we calculate the Euclidean distance between their descriptors. If the distance between two descriptors is small, we consider them as matching features.

```
1 def match_features(descriptors1, descriptors2):
2     """
3     Match features between two sets of descriptors using a brute-
4     force matcher.
5     """
6     bf = cv2.BFMatcher(cv2.NORM_L2, crossCheck=True)
7     matches = bf.match(descriptors1, descriptors2)
8     matches = sorted(matches, key=lambda x: x.distance)
9     return matches
```

## 6 Results

look at Figure 1 below, you will find that there are many arrows in Image3, but nearly no arrows and lines in other images. So we can say that our self-designed SIFT eligible for using.

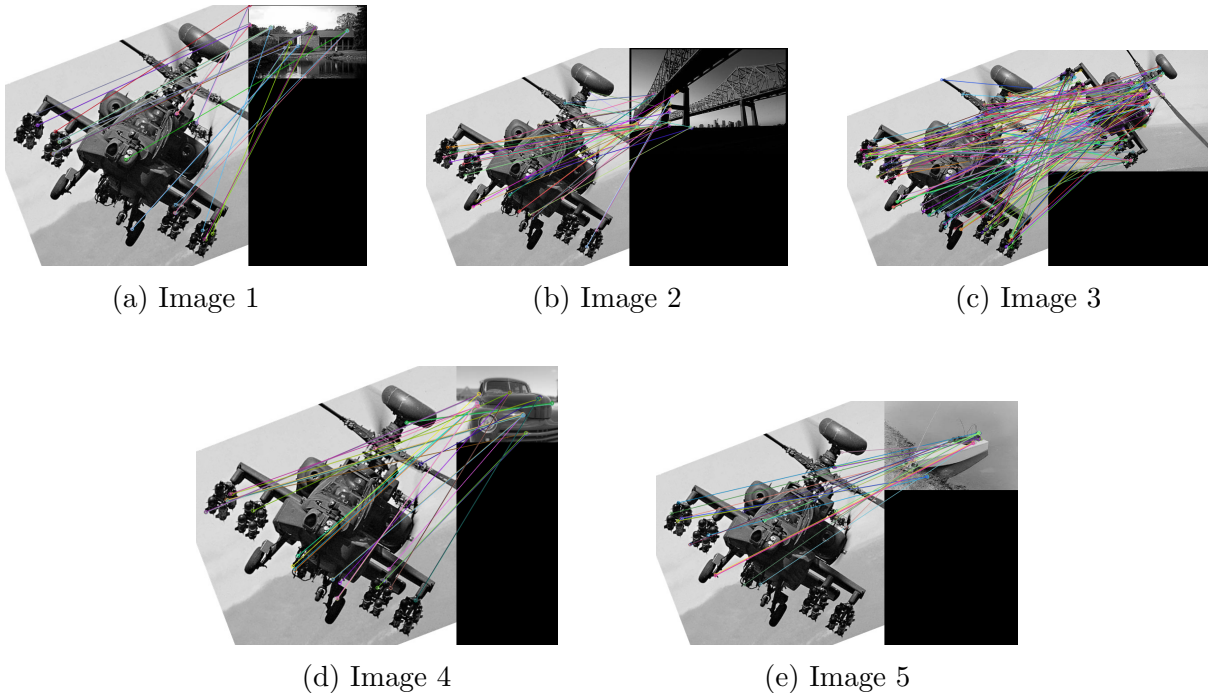




Figure 1: Result of Self Designed SIFT

Then we compare the self-designed result with the pre-built functions. Just look at Figure 2: The details are Shown in the table:



(a) self-designed SIFT



(b) pre-built SIFT

Figure 2: Comparison Of Different SIFT

Figure	self-designed	pre-built
1.jpg	32	31
2.jpg	54	35
3.jpg	241	399
4.jpg	33	31
5.jpg	26	18

We may not see the differences from Figure 2, but we can see the differences in the table above. It's obvious that the pre-built function is better than our self-designed function. From my perspective, the reason is that it's hard to adjust exactly the parameters of each function in our self-designed function. But for the pre-built function, maybe it has already used the best hyperparameters from the experiment.

## 7 Discussion

### 7.1 Why Delete Matched Points?

If matched points are not deleted, the same point may be matched to multiple points. Additionally, since the traversal is not area-based, points in different columns may remain.

### 7.2 Why Use Gaussian Blur?

Gaussian blur is applied to prevent sharp edges from causing distortion and inaccuracies in the results.

### 7.3 Challenges with Harris Corner Detection

Compared to direct SIFT, Harris corner detection produces many feature points, including many unnecessary ones. Using a Gaussian pyramid may mitigate this issue.

## 8 Reflections on the Lab

This lab provided significant opportunities for me to improve my skills in multiple areas, including mathematics, logical reasoning, and coding abilities. By implementing a simplified version of the SIFT algorithm, I gained a deeper understanding of mathematical concepts such as gradient calculations, Gaussian smoothing, and histogram interpolation.

From a logical thinking perspective, breaking down the SIFT descriptor into smaller steps, such as Gaussian blur, gradient computation, and descriptor construction, allowed me to approach the problem systematically.

Moreover, my coding abilities advanced significantly throughout the lab. Writing clean, efficient Python code while debugging errors and optimizing performance helped me build confidence in handling computer vision tasks.

Overall, this lab was a valuable experience that fostered growth in my mathematical understanding, logical thinking, and programming skills, laying a solid foundation for future exploration in computer vision and related fields.

## 9 Full code

pre built.py:

```
1 import cv2
2 import numpy as np
3 import os
4
5 def build_image_pyramid(image, levels=3):
6     """
7     Construct an image pyramid by resizing the image to multiple
8     scales.
9     """
10    pyramid = [image]
11    for _ in range(1, levels):
12        image = cv2.pyrDown(image) # Downscale the image
13        pyramid.append(image)
14    return pyramid
15
16 def compute_harris_corners(image, block_size=2, ksize=3, k=0.04):
17     """
18     Use Harris corner detection to find keypoints with an
19     adaptive threshold.
20     """
21    corners = cv2.cornerHarris(image, blockSize=block_size, ksize=ksize, k=k)
22    corners = cv2.dilate(corners, None) # Enhance corner points
23    adaptive_thresh = 0.25 * corners.max() # Compute adaptive
24    threshold based on maximum response
25    keypoints = np.argwhere(corners > adaptive_thresh)
26    keypoints = [cv2.KeyPoint(float(p[1]), float(p[0]), 1) for p
27                  in keypoints]
28    return keypoints
29
30 def compute_sift_descriptor(image, keypoints, patch_size=16):
31     """
32     Accurately compute SIFT-like descriptors with bilinear
33     interpolation for keypoints.
34     Each descriptor uses a 16x16 neighborhood divided into 4x4
```



```

31     cells,
    with an 8-bin histogram for each cell (128-dimensional
      descriptor).
32     """
33     descriptors = []
34     image = cv2.GaussianBlur(image, (5, 5), 1.6) # Gaussian
      smoothing to reduce noise.
35
36     for kp in keypoints:
37         x, y = int(kp.pt[0]), int(kp.pt[1]) # Keypoint
      coordinates.
38
39         # Extract the 16x16 neighborhood around the keypoint.
40         patch = image[max(0, y - patch_size // 2):y + patch_size
      // 2,
41                       max(0, x - patch_size // 2):x + patch_size
      // 2]
42
43         # Skip incomplete patches near image boundaries.
44         if patch.shape[0] != patch_size or patch.shape[1] !=
      patch_size:
45             continue
46
47         # Compute gradients using Sobel operators.
48         gx = cv2.Sobel(patch, cv2.CV_32F, 1, 0, ksize=3)
49         gy = cv2.Sobel(patch, cv2.CV_32F, 0, 1, ksize=3)
50         magnitude, angle = cv2.cartToPolar(gx, gy, angleInDegrees
      =True)
51
52         # Adjust angles relative to the dominant keypoint
      orientation.
53         main_orientation = kp.angle if kp.angle else 0
54         adjusted_angle = (angle - main_orientation) % 360
55
56         # Initialize a 4x4 grid with 8 bins per cell (128-
      dimensional descriptor).
57         cell_size = patch_size // 4 # Each cell is 4x4 pixels.
58         descriptor = np.zeros((4, 4, 8), dtype=np.float32)
59
60         # Iterate over all pixels in the 16x16 patch.
61         for i in range(patch_size):
62             for j in range(patch_size):
63                 # Compute relative coordinates within the grid.
64                 patch_x = j + 0.5 # Pixel center adjustment
65                 patch_y = i + 0.5
66
67                 cell_x = patch_x / cell_size # Grid row index
68                 cell_y = patch_y / cell_size # Grid column index
69
70                 # Identify the 4 nearest grid cells.
71                 x0, y0 = int(cell_x), int(cell_y)

```



```

72         x1, y1 = min(x0 + 1, 3), min(y0 + 1, 3) # Ensure
73             within 4x4 grid boundaries.
74
75         # Compute interpolation weights.
76         dx1, dy1 = cell_x - x0, cell_y - y0
77         dx2, dy2 = 1 - dx1, 1 - dy1
78
79         # Quantize angle into 8 bins (each 45 degrees).
80         bin_idx = int(adjusted_angle[i, j] // 45) % 8
81         magnitude_value = magnitude[i, j]
82
83         # Bilinearly interpolate the gradient magnitude
84         into the 4 nearest cells.
85         descriptor[y0, x0, bin_idx] += magnitude_value *
86             dx2 * dy2
87         descriptor[y0, x1, bin_idx] += magnitude_value *
88             dx1 * dy2
89         descriptor[y1, x0, bin_idx] += magnitude_value *
90             dx2 * dy1
91         descriptor[y1, x1, bin_idx] += magnitude_value *
92             dx1 * dy1
93
94         # Flatten the 4x4x8 grid into a 128-dimensional vector.
95         descriptor = descriptor.flatten()
96
97         # Normalize the descriptor to unit length (L2
98             normalization).
99         descriptor /= (np.linalg.norm(descriptor) + 1e-7)
100
101         # Clip values to 0.2 to suppress extreme gradients and re-
102             -normalize.
103         descriptor = np.clip(descriptor, 0, 0.2)
104         descriptor /= (np.linalg.norm(descriptor) + 1e-7)
105
106         descriptors.append(descriptor)
107
108     return np.array(descriptors, dtype=np.float32)
109
110
111 def match_features(descriptors1, descriptors2):
112     """
113     Match features between two sets of descriptors using a brute-
114         force matcher.
115     """
116     bf = cv2.BFMatcher(cv2.NORM_L2, crossCheck=True)
117     matches = bf.match(descriptors1, descriptors2)
118     matches = sorted(matches, key=lambda x: x.distance)
119     return matches
120
121
122 def main():
123     target_path = "target.jpg"

```

```

114     dataset_folder = "dataset"
115     output_folder = "output_matches"
116
117     # Ensure the output folder exists
118     if not os.path.exists(output_folder):
119         os.makedirs(output_folder)
120
121     # Load target image and create image pyramid
122     target_image = cv2.imread(target_path, cv2.IMREAD_GRAYSCALE)
123     target_pyramid = build_image_pyramid(target_image, levels=3)
124
125     # Compute Harris corners and SIFT descriptors for target
126     image
127     target_keypoints = compute_harris_corners(target_pyramid[0])
128     target_descriptors = compute_sift_descriptor(target_pyramid
129     [0], target_keypoints)
130
131     for image_name in os.listdir(dataset_folder):
132         image_path = os.path.join(dataset_folder, image_name)
133         dataset_image = cv2.imread(image_path, cv2.
134         IMREAD_GRAYSCALE)
135         dataset_pyramid = build_image_pyramid(dataset_image,
136         levels=3)
137
138         # Compute Harris corners and SIFT descriptors for dataset
139         image
140         dataset_keypoints = compute_harris_corners(
141         dataset_pyramid[0])
142         dataset_descriptors = compute_sift_descriptor(
143         dataset_pyramid[0], dataset_keypoints)
144
145         # Match features
146         matches = match_features(target_descriptors,
147         dataset_descriptors)
148
149         print(f"{image_name}: {len(matches)} matches")
150
151         # Draw matches
152         result_image = cv2.drawMatches(target_pyramid[0],
153         target_keypoints,
154         dataset_pyramid[0],
155         dataset_keypoints,
156         matches, None,
157         flags=cv2.
158         DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS
159         )
160
161         output_path = os.path.join(output_folder, f"matches_{
162         image_name}.png")

```

```

152         cv2.imwrite(output_path, result_image)
153
154 if __name__ == "__main__":
155     main()

```

self designed.py:

```

1  import cv2
2  import os
3  import numpy as np
4
5  # Path to the dataset folder and target image
6  dataset_path = "dataset"
7  target_image_path = "target.jpg"
8  output_folder = "system_shift"
9
10 # Ensure the output folder exists
11 if not os.path.exists(output_folder):
12     os.makedirs(output_folder)
13
14 # Load the target image
15 target_image = cv2.imread(target_image_path, cv2.IMREAD_GRAYSCALE)
16
17 if target_image is None:
18     print(f"Error: Could not load target image '{target_image_path}'")
19     exit()
20
21 # Initialize SIFT detector
22 sift = cv2.SIFT_create()
23
24 # Detect keypoints in the target image
25 target_keypoints = sift.detect(target_image, None)
26
27 # Compute descriptors for the detected keypoints
28 _, target_descriptors = sift.compute(target_image, target_keypoints)
29
30 # Draw keypoints for the target image (optional visualization)
31 target_keypoint_image = cv2.drawKeypoints(target_image, target_keypoints, None)
32 cv2.imwrite(os.path.join(output_folder, "target_keypoints.jpg"), target_keypoint_image)
33
34 # Set up FLANN-based matcher
35 FLANN_INDEX_KDTREE = 1
36 index_params = dict(algorithm=FLANN_INDEX_KDTREE, trees=5)
37 search_params = dict(checks=50)
38 flann = cv2.FlannBasedMatcher(index_params, search_params)
39
40 best_match_count = 0
41 best_match_image = None

```

```

41 best_match_keypoints = None
42 best_match_good_matches = None
43
44 # Iterate through all images in the dataset folder
45 for image_name in os.listdir(dataset_path):
46     image_path = os.path.join(dataset_path, image_name)
47
48     # Load the current image
49     image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
50     if image is None:
51         print(f"Warning: Could not load image '{image_path}'")
52         continue
53
54     # Detect keypoints in the current image
55     keypoints = sift.detect(image, None)
56
57     # Compute descriptors for the detected keypoints
58     _, descriptors = sift.compute(image, keypoints)
59
60     # Use FLANN matcher to find matches between the target and
61     current image
62     if descriptors is not None and target_descriptors is not None:
63         matches = flann.knnMatch(target_descriptors, descriptors,
64                                   k=2)
65
66         # Apply Lowe's ratio test to find good matches
67         good_matches = []
68         for m, n in matches:
69             if m.distance < 0.7 * n.distance:
70                 good_matches.append(m)
71
72         # Output the number of good matches for the current image
73         print(f"Image '{image_name}' has {len(good_matches)} good
74               matches.")
75
76         # Draw keypoints for the current image (optional
77         visualization)
78         keypoint_image = cv2.drawKeypoints(image, keypoints, None
79                                             )
80         cv2.imwrite(os.path.join(output_folder, f"keypoints_{
81                                   image_name}"), keypoint_image)
82
83         # Draw matches for the current image and save
84         result_image = cv2.drawMatches(target_image,
85                                         target_keypoints, image, keypoints, good_matches,
86                                         None, flags=cv2.
87                                             DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS
88                                         )
89         output_path = os.path.join(output_folder, f"matches_{
90                                   image_name}")

```

```

81         cv2.imwrite(output_path, result_image)
82
83         # Update the best match if the current image has more
84         good matches
85         if len(good_matches) > best_match_count:
86             best_match_count = len(good_matches)
87             best_match_image = image
88             best_match_keypoints = keypoints
89             best_match_good_matches = good_matches
90
91         # Draw and save the best match result if found
92         if best_match_image is not None:
93             best_match_result = cv2.drawMatches(target_image,
94                                                 target_keypoints, best_match_image, best_match_keypoints,
95                                                 best_match_good_matches,
96                                                 None, flags=cv2.
97                                                 DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
98
99             cv2.imwrite(os.path.join(output_folder, "best_match_result.
100                                jpg"), best_match_result)
101             print(f"Best match saved as 'best_match_result.jpg' with {
102                   best_match_count} good matches.")
103         else:
104             print("No good matches found in the dataset.")

```